


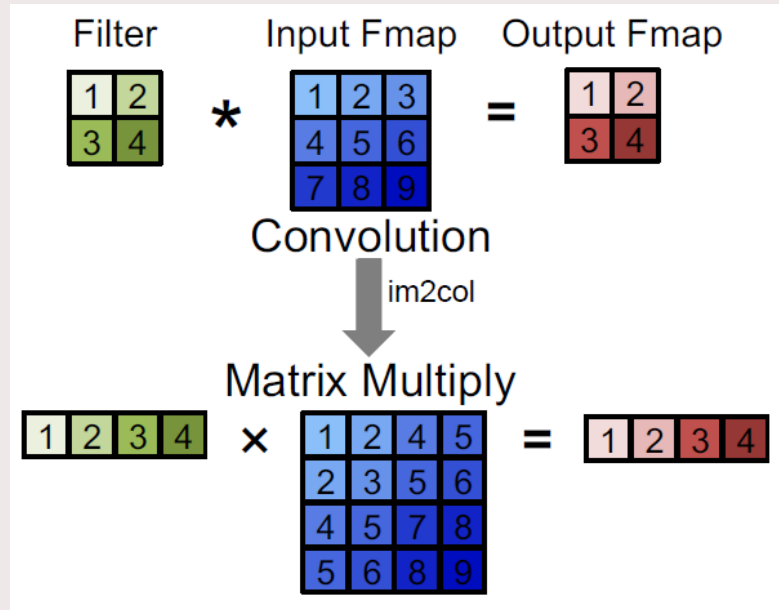


# Mapping and Accelerating a Convolutional Neural Network on a GPU




Haung Shao-Kai

# Part 1: Im2col



$$original\ memory = H * W * M * N * \frac{4}{1024} (KB)$$

$$im2col\ memory = R * S * (H - R + 1) * (W - S + 1) * M * N * \frac{4}{1024} (KB)$$




## Part 2: GPU

```
22: __global__ void optimized_fpu_gemm(fpu_type* A, fpu_type* B, float* C, float* D, int gemmM, int gemmN, int gemmK, float alpha, float beta){
23:     // implement the optimized floating point unit kernel using shared memory and loop unrolling
24:     // some step hints for you. Feel free to follow different steps
25:     // step 1. creat shared memory buffer
26:     __shared__ fpu_type A_tile[M_tiles_CUDA][K_tiles_CUDA];
27:     // __shared__ fpu_type B_tile[K_tiles_CUDA][M_tiles_CUDA]; //Uncoalesced
28:     __shared__ fpu_type B_tile[M_tiles_CUDA][K_tiles_CUDA]; //Coalesced
29:     // shorten parameters for clean re-use
30:     int tx = threadIdx.x;
31:     int ty = threadIdx.y;
32:     fpu_type accu = 0.0;
33:     // calculate current row and column of matrix C
34:     int row = blockIdx.y * blockDim.y + threadIdx.y;
35:     int col = blockIdx.x * blockDim.x + threadIdx.x;
36:     // step 2: main loop over K
37:     for (int tileIdx = 0; tileIdx < gemmK / K_tiles_CUDA; tileIdx++){
38:         // step 1: load data from global mem to shared mem
39:         if (tx + 1 <= K_tiles_CUDA && ty + 1 <= K_tiles_CUDA){
40:             A_tile[ty][tx] = A[row * gemmK + (tileIdx * K_tiles_CUDA + tx)];
41:             // B_tile[tx][ty] = B[col * gemmK + (tileIdx * K_tiles_CUDA + ty)]; //Uncoalesced
42:             B_tile[ty][tx] = B[col * gemmK + (tileIdx * K_tiles_CUDA + ty)]; //Coalesced
43:         }
44:         __syncthreads();
45:         // step 2: load data from shared mem to register
46:         for (int k = 0; k < K_tiles_CUDA; k++){
47:             // accu = accu + A_tile[ty][k] * B_tile[tx][k]; //Uncoalesced
48:             accu = accu + A_tile[ty][k] * B_tile[k][tx]; //Coalesced
49:         }
50:         __syncthreads();
51:     }
52:     // step 3: additional computations: adding matrix C
53:     accu = alpha * static_cast<float>(accu) + beta * C[row * gemmN + col];
54:     // step 4 store back the final results to globla memory
55:     D[row * gemmN + col] = accu;
56: }
```

## Baseline CUDA core

CPU_runtime	GPU_runtime	CNN_runtime	host_to_device_runtime
1.54E+06	1147.35	1647	31




```
1 Matrix Size : MxNxK : 4096x4096x4096
2 Warm-up Rounds = 0
3 TEST_ROUNDS = 1
4
5 average time of CPU baseline = 1.53906e+06 ms
6
7 Runtime of memory copying from host to device = 31 ms
8 average time of kernel simple_fpu = 1147.35 ms
9 Total runtime of the CNN layer = 1647 ms
```

$1540000 \div 1147.35 \approx 1342$  times faster

$1540000 \div 1647 \approx 935$  times faster

# Optimized CUDA core

Baseline_runtime	Uncoalesced_Optimized_runtime	Coalesced_Optimized_runtime
1147.35	406.769	141.204



## Uncoalesced tiling

```
1 Matrix Size : MxNxK : 4096x4096x4096
2 numeric check passed
3 Tiling size = 32
4 average time of kernel optimized_fpu = 406.769 ms
```

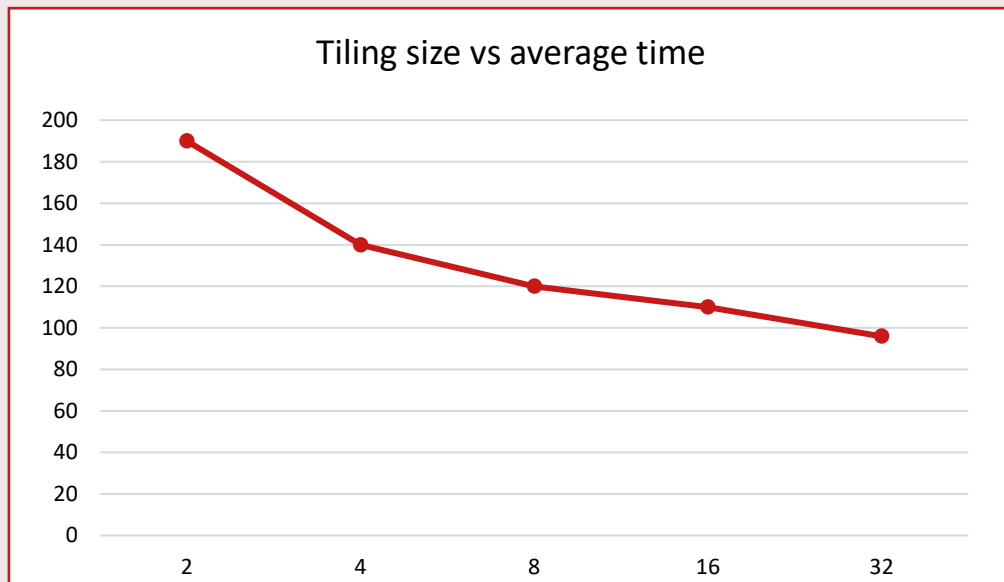
## Coalesced tiling

```
1 Matrix Size : MxNxK : 4096x4096x4096
2 Runtime of matrix transpose = 376 ms
3 numeric check passed
4 Tiling size = 32
5 average time of kernel optimized_fpu = 141.204 ms
```

$1147.35 \div 406.769 \approx 3 \text{ times faster}$

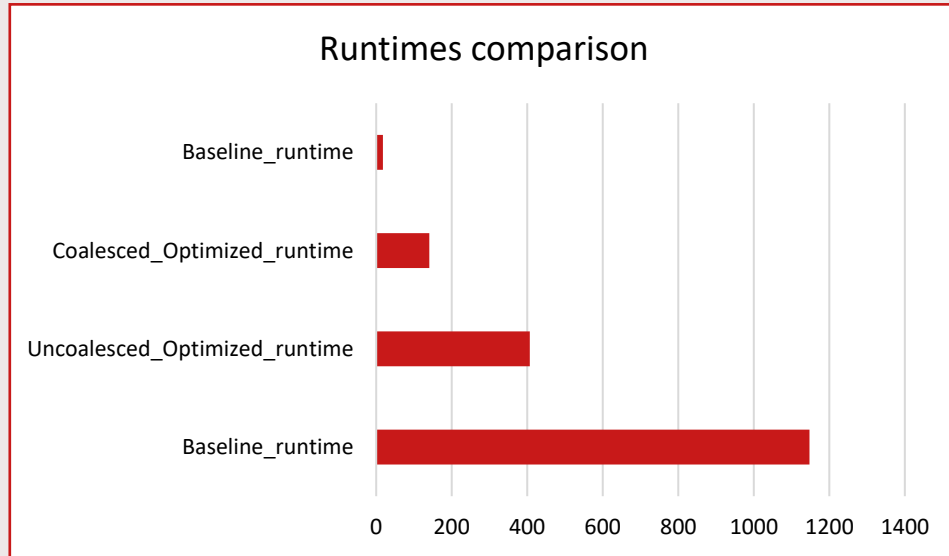
$1147.35 \div 141.204 \approx 8 \text{ times faster}$

## Tiling size vs Average time



## Baseline Tensor core

CUDA			Tensor
Baseline_runtime	Uncoalesced_Optimized_runtime	Coalesced_Optimized_runtime	Baseline_runtime
1147.35	406.769	141.204	17.9



average time of kernel simple\_wmma = 17.9091 ms  
Total runtime of the CNN layer = 363 ms

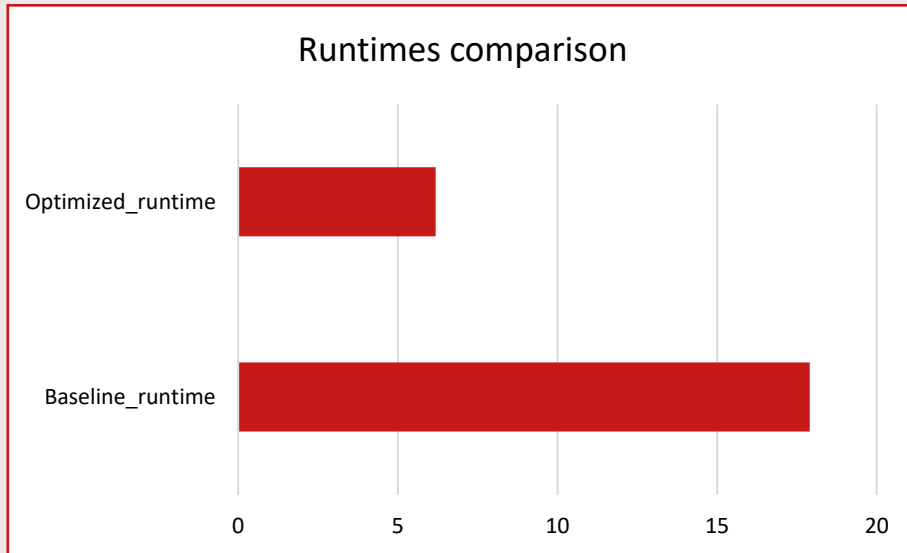
$1147.35 \div 17.9 \approx 64 \text{ times faster}$

$406.769 \div 17.9 \approx 23 \text{ times faster}$

$141.204 \div 17.9 \approx 8 \text{ times faster}$

## Optimized Tensor core (shared memory only)

Tensor	
Baseline_runtime	Optimized_runtime
17.9	6.18




```
average time of kernel optimized_wmma = 6.18253 ms  
Total runtime of the CNN layer = 352 ms
```

$17.9 \div 6.18 \approx 3 \text{ times faster}$



# Comparison of total run time

Total Runtime				
CPU	simple_fpu	optimized_fpu	simple_wmma	optimized_wmma
1.54E+06	1647	833	363	352
Kernel Runtime				
1.54E+06	1147.35	111.72	17.9	6.18



```

1  Matrix Size : MxNxK : 4096x4096x4096
2  Warm-up Rounds = 0
3  TEST_ROUNDS = 1
4
5  average time of CPU baseline = 1.53906e+06 ms
6
7  Runtime of memory copying from host to device = 31 ms
8  average time of kernel simple_fpu = 1147.35 ms
9  Total runtime of the CNN layer = 1647 ms
10
11 average time of kernel simple_wmma = 17.9091 ms
12 Total runtime of the CNN layer = 363 ms
13
14 Runtime of matrix transpose = 376 ms
15 Tiling size = 32
16 average time of kernel optimized_fpu = 111.721 ms
17 Total runtime of the CNN layer = 833 ms
18
19 average time of kernel optimized_wmma = 6.18253 ms
20 Total runtime of the CNN layer = 352 ms

```